




Transform Array Operations to Optimize Memory Access for GPGPU

Mads Ynddal
SJT402



Overview

- | | |
|--------------------|----------------|
| 1. Motivation | 4. Benchmarks |
| 2. Access Patterns | 5. Future Work |
| 3. Reduction | 6. Conclusion |

I. Motivation

107

Bh

Bohrium

[270]

Motivation

- Bohrium
 - Novice
 - Experts
- Performance Gap
 - Reductions
 - Memory Throughput



Motivation

- Bohrium
- Novice
- Experts
- Performance Gap
 - Reductions
 - Memory Throughput

NumPy Stencil

```
center = grid[1:-1, 1:-1]
north = grid[0:-2, 1:-1]
east = grid[1:-1, 2:]
west = grid[1:-1, 0:-2]
south = grid[2:, 1:-1]
work = 0.2 * (center + north +
              east + west + south)
delta = np.sum(np.abs(work-center))
center[:] = work
```

Motivation


- Bohrium
- Novice
- Experts
- Performance Gap
- Reductions
- Memory Throughput

OpenCL Kernel Stencil

```
uint gid = get_global_id(0);
if (gid >= width)
    return;
uint w = width + 2;
uint h = height + 2;
DTYPE d = 0.0;
for (uint i = 0; i < height; ++i)
{
    uint offset = i*w;
    DTYPE up      = in[gid+1+offset];
    DTYPE left    = in[gid+w+offset];
    DTYPE right   = in[gid+w+2+offset];
    DTYPE down    = in[gid+1+w*2+offset];
    DTYPE center  = in[gid+w+1+offset];
    DTYPE out_center =
        (center + up + left + right + down) * 0.2;
    out[gid+w+1+offset] = out_center;
    d += fabs(out_center - center);
}
delta[gid] = d;
```



Motivation

- Bohrium
 - Novice
 - Experts
- Performance Gap
- Reductions
- Memory Throughput



Motivation

- Bohrium
 - Novice
 - Experts
- Performance Gap
 - Reductions
 - Memory Throughput



Motivation


- Bohrium
 - Novice
 - Experts
- Performance Gap
 - Reductions
- Memory Throughput

Nvidia Guidelines

“Bandwidth is one of the most important gating factors for performance. Almost all changes to code should be made in the context of how they affect bandwidth”


2. Access Patterns

- Tensors
- Inspection
- Changing Access Pattern




Tensors

- Rank
- Strides
- Operations



Tensors


- Rank
- Strides
- Operations



Inspection

Rank-4 tensor with dimensions:
 $50 \times 50 \times 50 \times 50$

- Tensor Example
- Generated Code
- Current Pattern
- Coalesced Pattern



Inspection

Rank-4 tensor with dimensions:
 $50 \times 50 \times 50 \times 50$

- Tensor Example
- Generated Code
- Current Pattern
- Coalesced Pattern

```
for (ulong i0 = 0; i0 < 50; ++i0) {  
    for (ulong i1 = 0; i1 < 50; ++i1) {  
        for (ulong i2 = 0; i2 < 50; ++i2) {  
            for (ulong i3 = 0; i3 < 50; ++i3) {  
                a0[+i0*125000 + i1*2500 + i2*50 + i3] = 1;  
            }  
        }  
    }  
}
```

Inspection

Rank-4 tensor with dimensions:
 $50 \times 50 \times 50 \times 50$

- Tensor Example
- Generated Code
- Current Pattern
- Coalesced Pattern

```
const uint g0 = get_global_id(0);  
if (g0 >= 50) { return; } // Prevent overflow  
  
{const ulong i0 = g0;  
  for (ulong i1 = 0; i1 < 50; ++i1) {  
    for (ulong i2 = 0; i2 < 50; ++i2) {  
      for (ulong i3 = 0; i3 < 50; ++i3) {  
        a0[+i0*125000 + i1*2500 + i2*50 + i3] = 1;  
      }  
    }  
  }  
}
```


Inspection

Rank-4 tensor with dimensions:
 $50 \times 50 \times 50 \times 50$

- Tensor Example
- Generated Code
- Current Pattern
- Coalesced Pattern

```
const uint g3 = get_global_id(0);  
if (g3 >= 50) { return; } // Prevent overflow  
  
for (ulong i0 = 0; i0 < 50; ++i0) {  
    for (ulong i1 = 0; i1 < 50; ++i1) {  
        for (ulong i2 = 0; i2 < 50; ++i2) {  
            {const ulong i3 = g3;  
                a0[+i0*125000 + i1*2500 + i2*50 + i3] = 1;  
            }  
        }  
    }  
}
```


Changing Access Pattern

- Decide on Parallel Axis
- Implementation
- Results

Truncated Block Structure of Kernel

```
rank: 0, size: 75, sweeps: {}  
  rank: 1, size: 75, sweeps: {}  
    rank: 2, size: 75, sweeps: {  
      BH_ADD_REDUCE a0 a1 3}  
    rank: 3, size: 75, sweeps: {  
      BH_ADD_REDUCE a2 a3 4}
```

Changing Access Pattern


- Decide on Parallel Axis
- Implementation
- Results

Remapped Iterators

```
uint g2 = get_global_id(0);  
uint g1 = get_global_id(1);  
uint g0 = get_global_id(2);  
// ...  
i0 = g0  
i1 = g1  
i2 = g2
```


Changing Access Pattern

- Decide on Parallel Axis
- Implementation
- Results




3. Reduction

- Hardware Constraints
- Access Pattern in Reductions



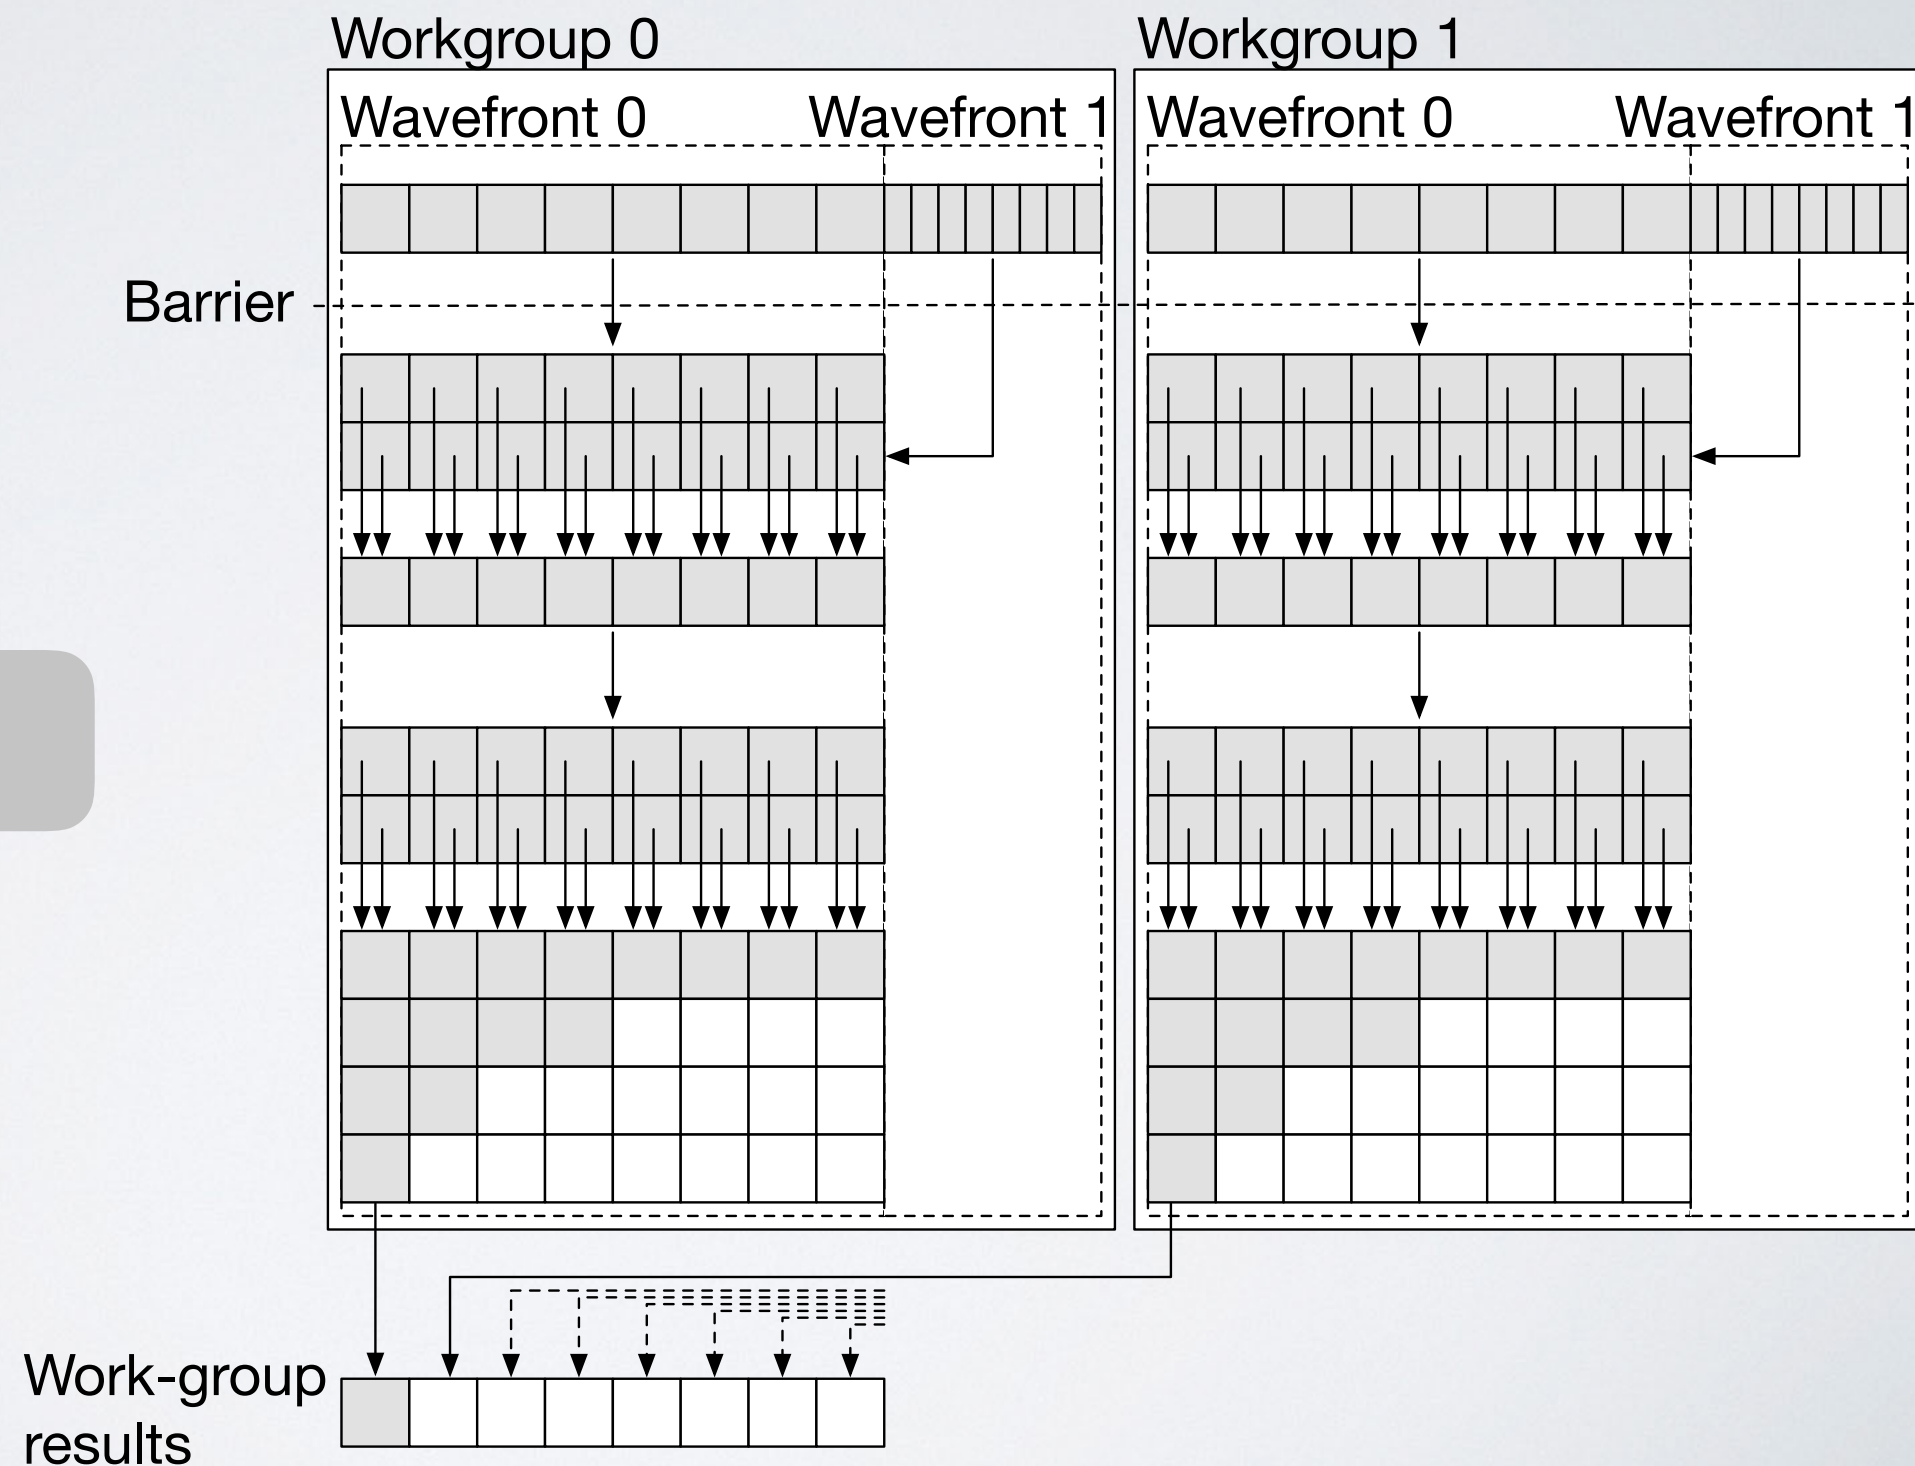
Hardware Constraints

- Wavefronts
- Work-groups
- Bohrium



Hardware Constraints


- Wavefronts
- Work-groups
- Bohrium




Hardware Constraints

- Wavefronts
- Work-groups
- Bohrium

Old vector reductions




New vector reductions



Access Pattern in Reductions


- Axis
- Column-wise
- Row-wise
- Transposing



0,0,0				0,1,0				0,2,0			0,2,3
1,0,0				1,1,0				1,2,0			1,2,3


Access Pattern in Reductions

- Axis
- Column-wise
- Row-wise
- Transposing



Access Pattern in Reductions


- Axis
- Column-wise
- Row-wise
- Transposing



Segmented Reduction Preferred


Access Pattern in Reductions

- Axis
- Column-wise
- Row-wise
- Transposing



4. Benchmarks

- Benchmark Suite
- Handwritten vs. Bohrium

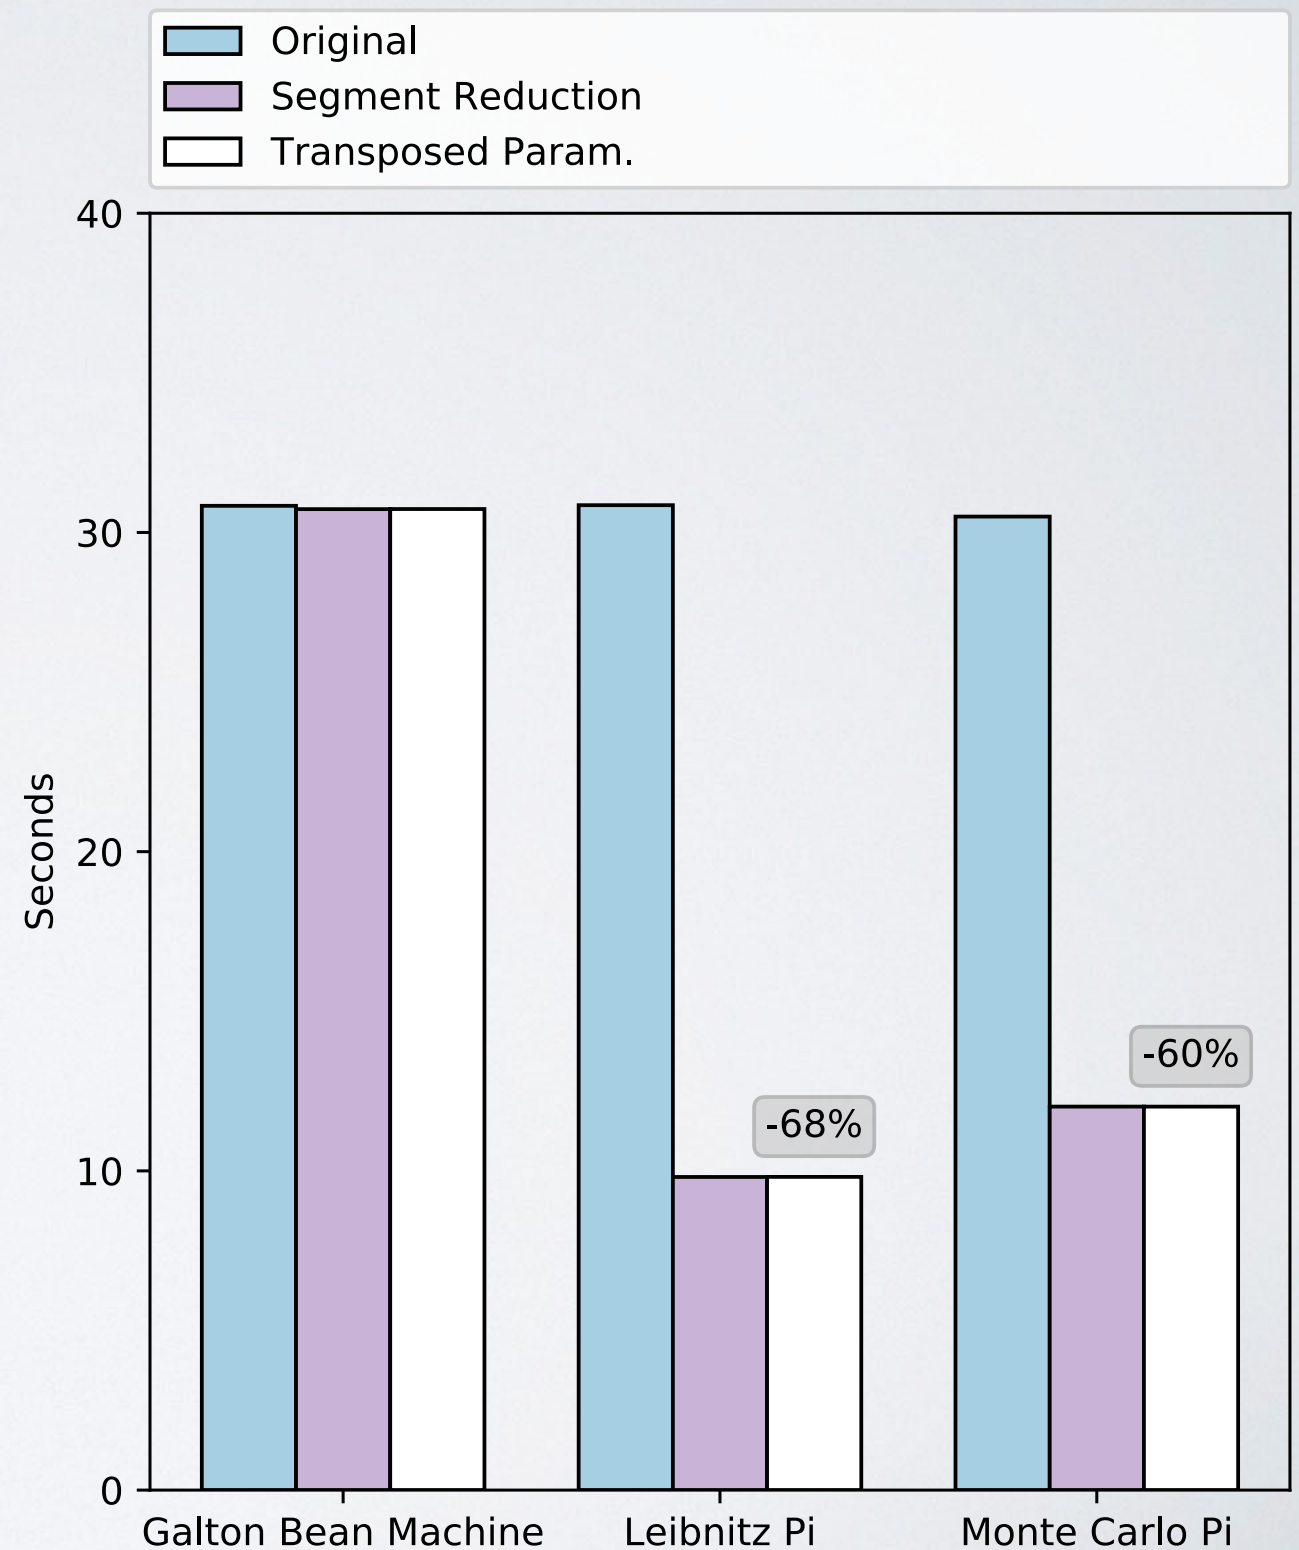


Benchmark Suite

- Embarrassingly Parallel
- Finite Difference
- Streaming/Reduction

Benchmark Suite

- Embarrassingly Parallel
 - I. Galton Bean Machine
 - II. Leibnitz Pi
 - III. Monte Carlo Pi
- Finite Difference
- Streaming/Reduction



Benchmark Suite

- Embarrassingly Parallel


- Finite Difference

I. Heat Equation

II. Lattice Boltzmann

III. Shallow Water

- Streaming/Reduction




Benchmark Suite

- Embarrassingly Parallel
- Finite Difference
- Streaming/Reduction

I. Magnetic Field Extrapolation


II. N-body

III. X-ray Sim



Handwritten vs. Bohrium

- Handwritten version
- NumPy version
- Complexity of simulations



5. Future Work

- Problem with MFE
- Throughput
- Next Focus Points

v2.0

Problem with MFE

- Access Pattern
- Handling the Access Pattern

MFE Kernel Stripped of Constant Strides
Param. $128 \times 1 \times 1$

```
s8_8 = a8[ +i3*75 +i4];  
t6 = a7[ +i2*5625 +i3*75 +i4]  
      * s8_8;
```

MFE Kernel Stripped of Constant Strides
Param. $1 \times 2 \times 64$

```
s8_8 = a8[ +i0*5625 +i3*75 +i4];  
t6 = a7[ +i3*75 +i4] * s8_8;
```

Problem with MFE

- Access Pattern

- Handling the Access Pattern

MFE NumPy Code

```
np.sum(  
    np.sum(  
        temp_x * exprx[:, None, None],  
        -1),  
    -1)
```

temp_x exprx[:,None,None]
(75,75,75,75) * (75,1,1,75,75)




2x
reduce (75,75,75,75,75)



(75,75,75)


Throughput

- Memory Copy
- Vector Reduction
- Segmented Reduction



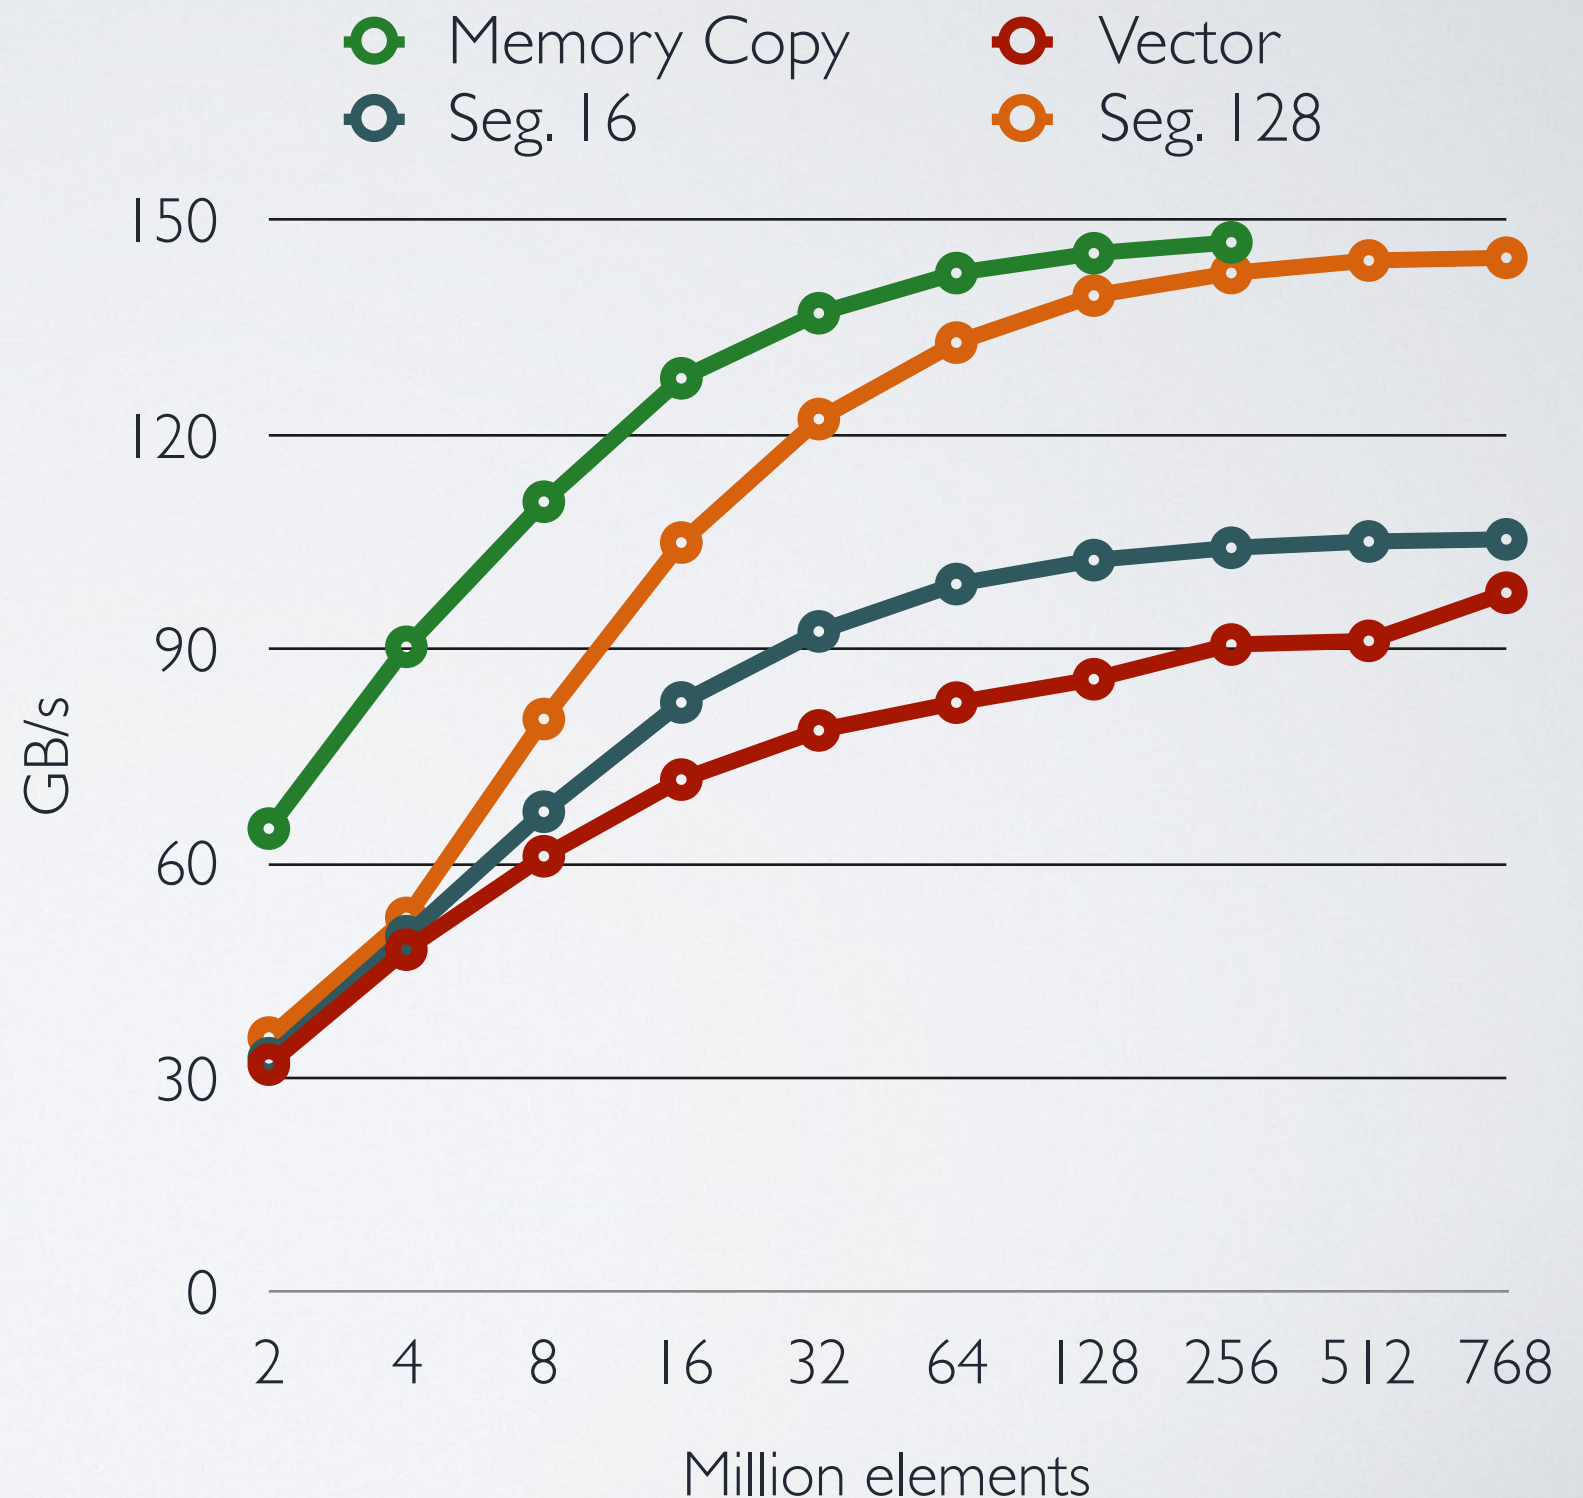
Throughput

- Memory Copy
- Vector Reduction
- Segmented Reduction



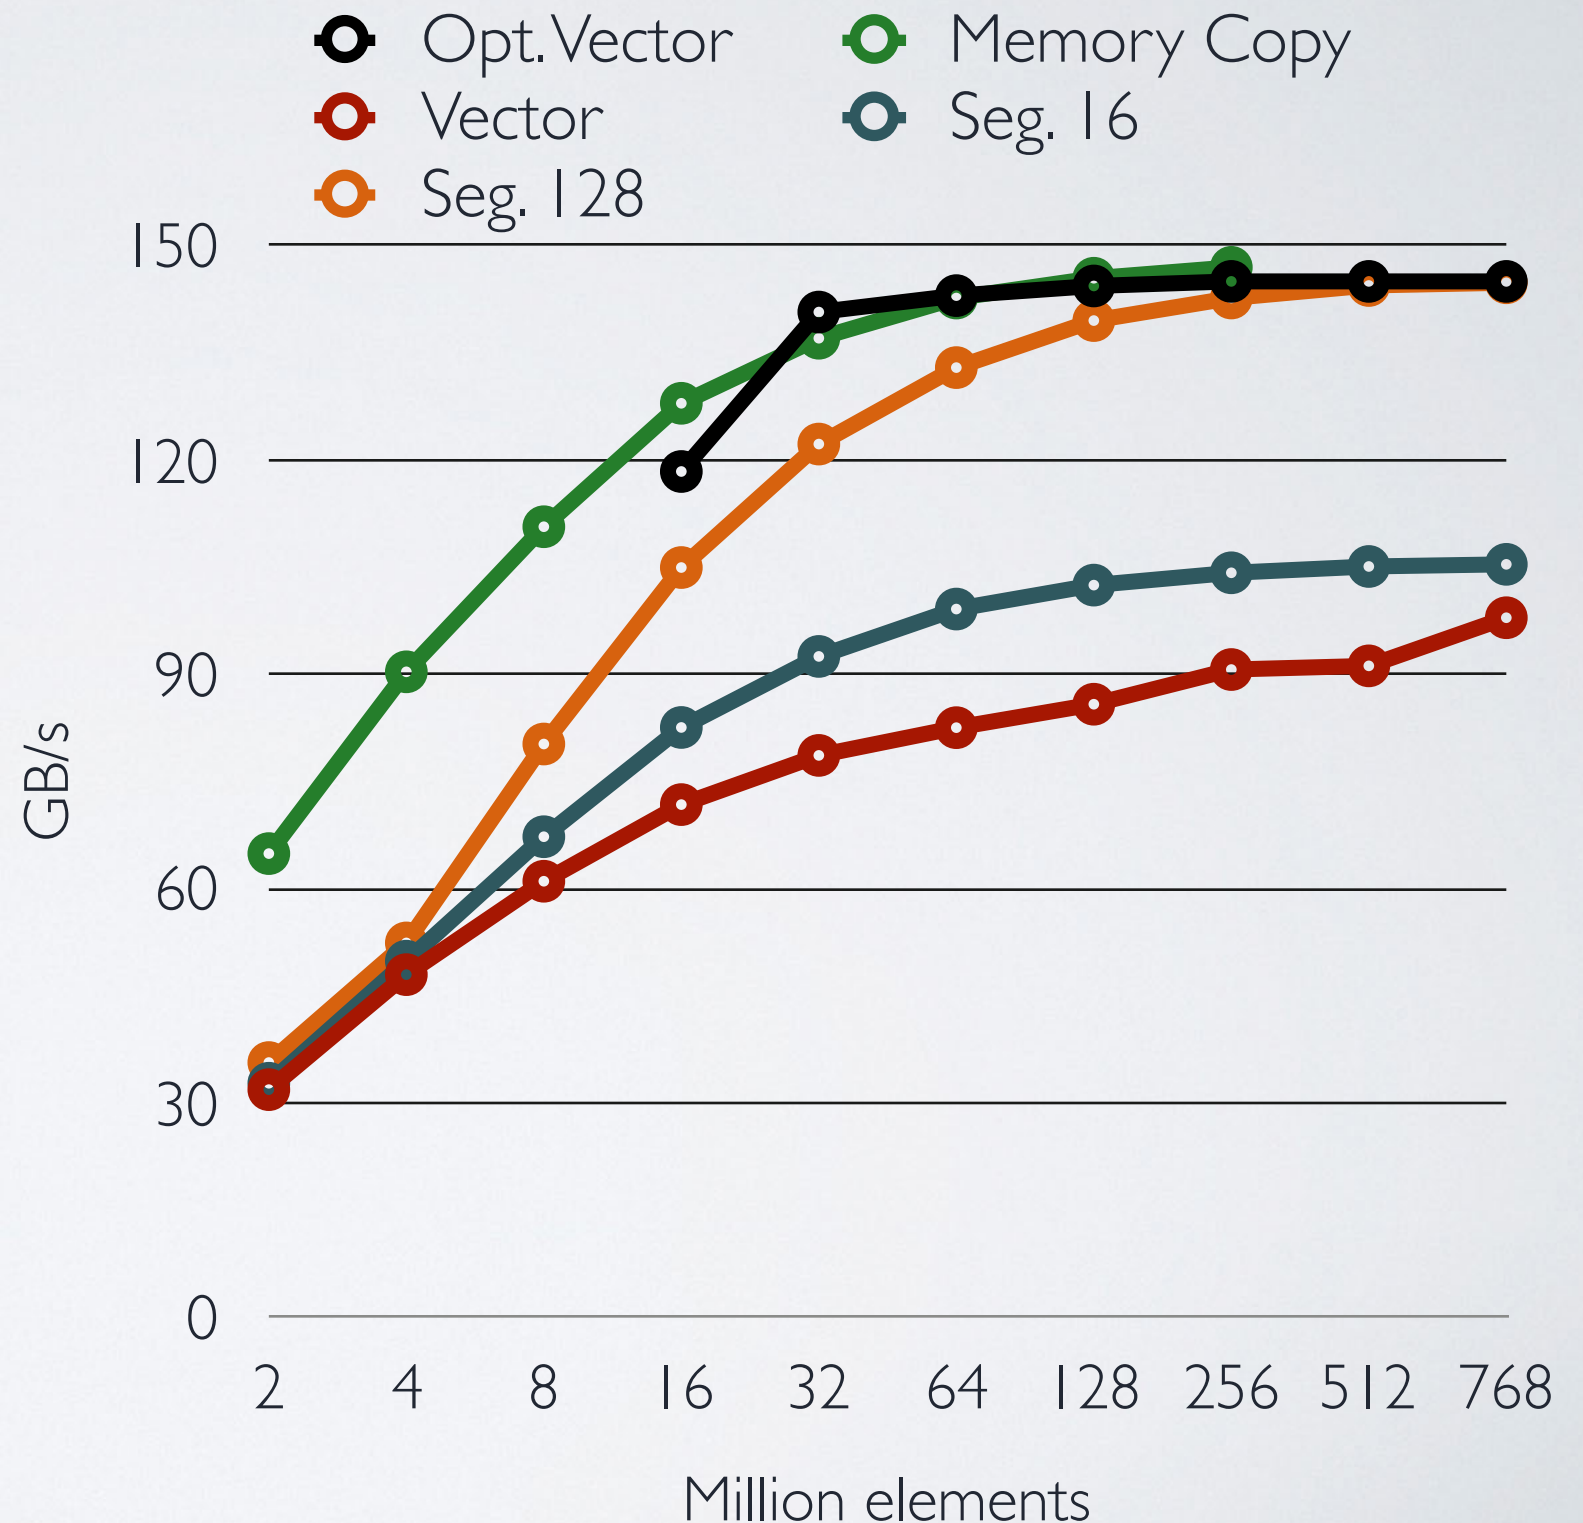
Throughput

- Memory Copy
- Vector Reduction
- Segmented Reduction




Throughput

- Memory Copy
- Vector Reduction
- Segmented Reduction



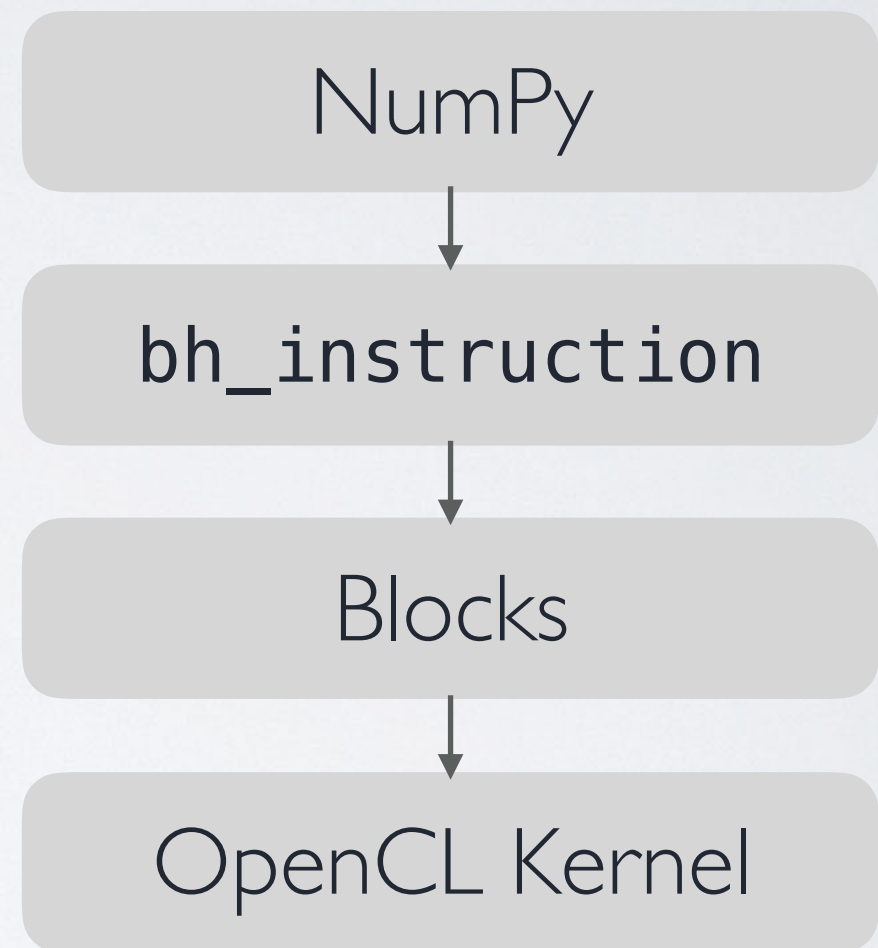
Next Focus Points

- Memory Throughput
- “Chicken or the egg” Issues
- Kernel-based Parameters
- Reshape and Transform



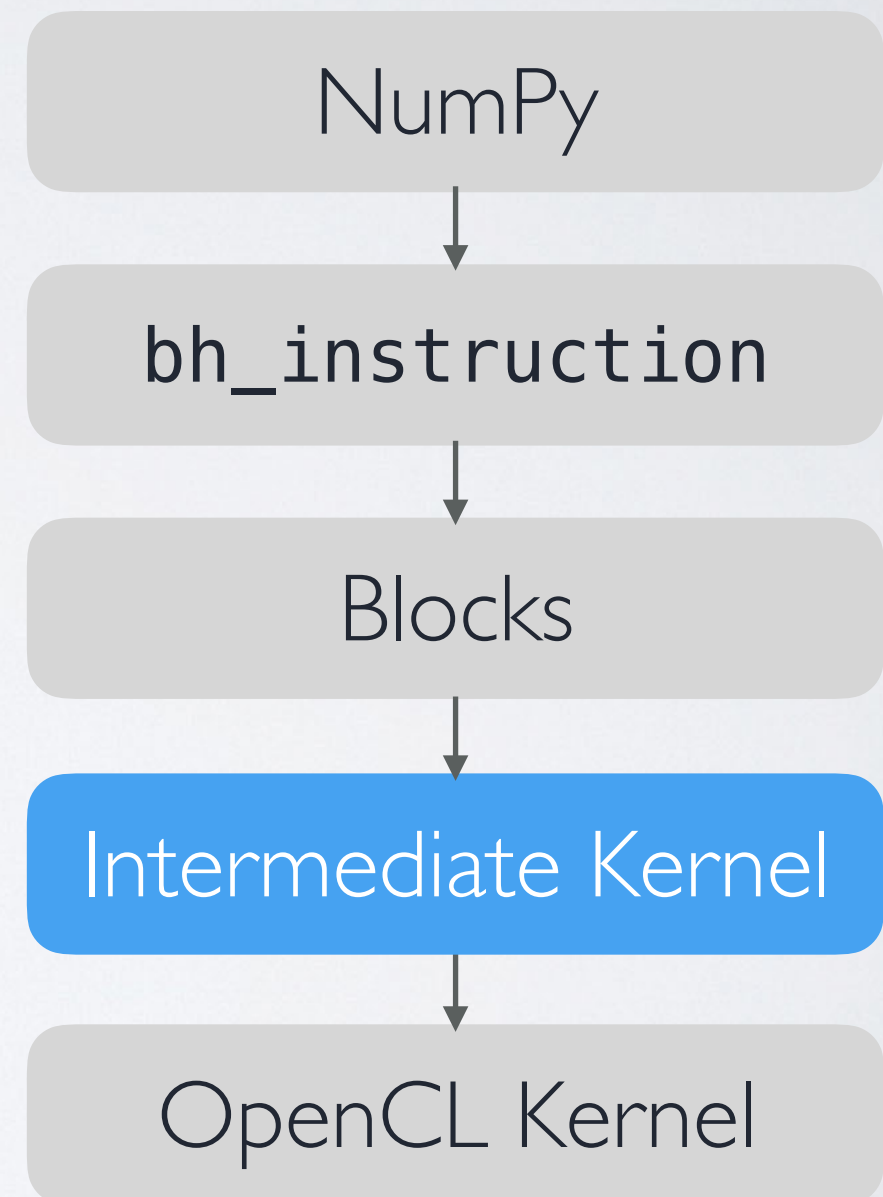
Next Focus Points

- Memory Throughput
- “Chicken or the egg” Issues
- Kernel-based Parameters
- Reshape and Transform



Next Focus Points

- Memory Throughput
- “Chicken or the egg” Issues
- Kernel-based Parameters
- Reshape and Transform



Next Focus Points

- Memory Throughput
- “Chicken or the egg” Issues
- Kernel-based Parameters
- Reshape and Transform


```
exprx[:,None,None]  
(75,1,1,75,75)
```



```
uint g2 = get_global_id(0);  
uint g1 = get_global_id(1);  
uint g0 = get_global_id(2);
```

Next Focus Points

- Memory Throughput
- “Chicken or the egg” Issues
- Kernel-based Parameters
- Reshape and Transform



6. Conclusion

- Vector Reduction
- Access Pattern
- Performance Gap
- Reflection

6. Conclusion

- Vector Reduction
- Access Pattern
- Performance Gap
- Reflection

6. Conclusion

- Vector Reduction
- Access Pattern
- Performance Gap
- Reflection

6. Conclusion

- Vector Reduction
- Access Pattern
- Performance Gap
- Reflection